

ML_Lab_3 (1)

December 3, 2023

```
[1]: import h5py
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Model, Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, Add,
↳Activation
from sklearn.metrics import accuracy_score
import keras
```

```
/home/as16494/.local/lib/python3.8/site-packages/numpy/core/getlimits.py:518:
UserWarning: The value of the smallest subnormal for <class 'numpy.float32'>
type is zero.
```

```
    setattr(self, word, getattr(machar, word).flat[0])
```

```
/home/as16494/.local/lib/python3.8/site-packages/numpy/core/getlimits.py:89:
UserWarning: The value of the smallest subnormal for <class 'numpy.float32'>
type is zero.
```

```
    return self._float_to_str(self.smallest_subnormal)
```

```
/home/as16494/.local/lib/python3.8/site-packages/numpy/core/getlimits.py:518:
UserWarning: The value of the smallest subnormal for <class 'numpy.float64'>
type is zero.
```

```
    setattr(self, word, getattr(machar, word).flat[0])
```

```
/home/as16494/.local/lib/python3.8/site-packages/numpy/core/getlimits.py:89:
UserWarning: The value of the smallest subnormal for <class 'numpy.float64'>
type is zero.
```

```
    return self._float_to_str(self.smallest_subnormal)
```

```
[2]: import tensorflow as tf
badnet_model = tf.keras.models.load_model('bd_net.h5')
```

```
[3]: badnet_model.summary()
```

Model: "model_1"

```
-----
-----
Layer (type)                Output Shape              Param #   Connected to
=====
input (InputLayer)          [(None, 55, 47, 3)]      0         []
```

conv_1 (Conv2D) ['input[0][0]']	(None, 52, 44, 20)	980
pool_1 (MaxPooling2D) ['conv_1[0][0]']	(None, 26, 22, 20)	0
conv_2 (Conv2D) ['pool_1[0][0]']	(None, 24, 20, 40)	7240
pool_2 (MaxPooling2D) ['conv_2[0][0]']	(None, 12, 10, 40)	0
conv_3 (Conv2D) ['pool_2[0][0]']	(None, 10, 8, 60)	21660
pool_3 (MaxPooling2D) ['conv_3[0][0]']	(None, 5, 4, 60)	0
conv_4 (Conv2D) ['pool_3[0][0]']	(None, 4, 3, 80)	19280
flatten_1 (Flatten) ['pool_3[0][0]']	(None, 1200)	0
flatten_2 (Flatten) ['conv_4[0][0]']	(None, 960)	0
fc_1 (Dense) ['flatten_1[0][0]']	(None, 160)	192160
fc_2 (Dense) ['flatten_2[0][0]']	(None, 160)	153760
add_1 (Add) ['fc_1[0][0]', 'fc_2[0][0]']	(None, 160)	0
activation_1 (Activation) ['add_1[0][0]']	(None, 160)	0
output (Dense) ['activation_1[0][0]']	(None, 1283)	206563

=====

Total params: 601643 (2.30 MB)
Trainable params: 601643 (2.30 MB)

Non-trainable params: 0 (0.00 Byte)

```
[4]: def load_dataset(h5_file_path):  
    with h5py.File(h5_file_path, 'r') as file:  
        features = np.array(file['data'])  
        labels = np.array(file['label'])  
        features = np.transpose(features, (0, 2, 3, 1))  
  
    return features, labels
```

```
[5]: cl_x_valid, cl_y_valid = load_dataset('valid.h5')  
    cl_x_test, cl_y_test = load_dataset('test.h5')  
    bd_x_valid, bd_y_valid = load_dataset('bd_valid.h5')  
    bd_x_test, bd_y_test = load_dataset('bd_test.h5')  
    N = int(cl_y_test.max())
```

```
[6]: backdoor_model = keras.models.load_model('bd_net.h5')  
  
def evaluate_model_performance(model, clean_data, clean_labels, backdoor_data,  
    ↪ backdoor_labels):  
    clean_predictions = np.argmax(model.predict(clean_data), axis=1)  
    clean_accuracy = accuracy_score(clean_labels, clean_predictions) * 100  
    backdoor_predictions = np.argmax(model.predict(backdoor_data), axis=1)  
    attack_success_rate = accuracy_score(backdoor_labels, backdoor_predictions) ↪  
    ↪ * 100  
    return clean_accuracy, attack_success_rate  
  
# Evaluate the model  
clean_accuracy, attack_success_rate = evaluate_model_performance(  
    backdoor_model, cl_x_test, cl_y_test, bd_x_test, bd_y_test  
)  
  
print('Clean Classification Accuracy:', clean_accuracy)  
print('Attack Success Rate:', attack_success_rate)
```

```
401/401 [=====] - 2s 3ms/step  
401/401 [=====] - 1s 3ms/step  
Clean Classification Accuracy: 98.62042088854248  
Attack Success Rate: 100.0
```

```
[7]: def createPrunedModel(original_model, X_threshold):  
    # Copy the model structure and weights  
    modified_model = keras.models.clone_model(original_model)  
    modified_model.set_weights(original_model.get_weights())
```

```

initial_predictions = np.argmax(original_model.predict(cl_x_valid), axis=1)
initial_accuracy = np.mean(initial_predictions == cl_y_valid) * 100

# Target the specific layer for pruning
target_layer = modified_model.get_layer('conv_3')
activation_model = keras.Model(inputs=modified_model.input,
→outputs=target_layer.output)
channel_activations = activation_model.predict(cl_x_valid).sum(axis=(0, 1,
→2))

channels_sorted_by_activation = np.argsort(channel_activations)

for channel_index in channels_sorted_by_activation:
    # Modify the weights of the target layer to "remove" a channel
    layer_weights = target_layer.get_weights()
    layer_weights[0][:, :, :, channel_index] = 0 # Set the weights of the
→channel to zero
    target_layer.set_weights(layer_weights)

    modified_predictions = np.argmax(modified_model.predict(cl_x_valid),
→axis=1)
    modified_accuracy = np.mean(modified_predictions == cl_y_valid) * 100

    if initial_accuracy - modified_accuracy > X_threshold:
        target_layer.set_weights(target_layer.get_weights())
        break

return modified_model

```

```

[8]: def evaluateModelPerformance(pruned_model, original_model):
    # Accuracy evaluation on clean test data
    pred_pruned_clean = np.argmax(pruned_model.predict(cl_x_test), axis=1)
    pred_original_clean = np.argmax(original_model.predict(cl_x_test), axis=1)

    clean_predictions = [pred if pred == pred_original else N + 1 for pred,
→pred_original in zip(pred_pruned_clean, pred_original_clean)]
    accuracy = np.mean(np.array(clean_predictions) == cl_y_test) * 100

    # ASR evaluation on backdoored test data
    pred_pruned_bd = np.argmax(pruned_model.predict(bd_x_test), axis=1)
    pred_original_bd = np.argmax(original_model.predict(bd_x_test), axis=1)

    bd_predictions = [pred if pred == pred_original else N + 1 for pred,
→pred_original in zip(pred_pruned_bd, pred_original_bd)]
    asr = np.mean(np.array(bd_predictions) == bd_y_test) * 100

    return accuracy, asr

```

```
[9]: Xs = [2, 4, 10]
    for X in Xs:

        modified_model = createPrunedModel(badnet_model, X)
        # Assess the performance of the pruned model
        acc, asr = evaluateModelPerformance(modified_model, badnet_model)
        print(f"Drop = {X}%\n\tAccuracy of Pruned Model = {acc}%\n\tAttack Success_
↪Rate = {asr}%\n")
        modified_model.save(f"models/bd_prime_{X}.h5")
```

[illegible]

```

361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
401/401 [=====] - 1s 3ms/step
401/401 [=====] - 1s 3ms/step
401/401 [=====] - 1s 3ms/step
401/401 [=====] - 1s 3ms/step

```

Accuracy of Pruned Model = 95.8846453624318%

```
WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet
to be built. `model.compile_metrics` will be empty until you train or evaluate
the model.
```

[illegible]

[illegible]

Drop = 4%

Accuracy of Pruned Model = 94.61418550272798%

Attack Success Rate = 99.97661730319564%

```
WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet
to be built. `model.compile_metrics` will be empty until you train or evaluate
the model.
```

```
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 2ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
361/361 [=====] - 1s 3ms/step
```

[illegible]


```
401/401 [=====] - 1s 3ms/step
401/401 [=====] - 1s 3ms/step
Drop = 10%
    Accuracy of Pruned Model = 84.45830085736556%
    Attack Success Rate = 76.1730319563523%
```

```
WARNING:tensorflow:Compiled the loaded model, but the compiled metrics have yet
to be built. `model.compile_metrics` will be empty until you train or evaluate
the model.
```

```
[ ]:
```